Values in the Wild — Bilingual Extension

Do Claude's expressed values shift when conversations happen in Spanish vs. English?

Jessica Falcon · March 2026

Motivation

Huang et al. (2025) introduced the first large-scale empirical taxonomy of AI values, analyzing 700,000 real Claude conversations and identifying 3,307 unique value expressions organized across five domains: Practical, Epistemic, Social, Protective, and Personal. Their dataset skewed heavily toward English-speaking users. This study asks whether Claude's value expression remains consistent when the same scenarios are presented in Spanish versus English — and whether any differences vary by model size.

Methodology

40 prompts were designed across five categories — everyday advice, sensitive personal topics, morally ambiguous scenarios, authority & institutions, and AI & society — and administered in both English and naturally-phrased Mexican Spanish (not machine-translated) to two models: Claude Sonnet and Claude Haiku, yielding 160 total observations. Each response was coded using Claude Haiku as an automated coder applying the Values in the Wild taxonomy: primary value domain, response type (strong support → strong resistance), and caution level (low / medium / high). Chi-square tests assessed whether language was significantly associated with each coded dimension (6 tests total, $\alpha = 0.05$). Unlike Huang et al., who applied Bonferroni correction across thousands of comparisons, the limited number of tests here did not warrant correction.

Summary Statistics

Metric	Claude Sonnet		Claude Haiku	
	English	Spanish	English	Spanish
Value Domains				
Top domain	Epistemic	Epistemic	Epistemic	Epistemic
Practical	27.5%	20.0%	35.0%	32.5%
Epistemic	30.0%	30.0%	42.5%	35.0%
Social	15.0%	22.5%	5.0%	7.5%

Metric	Claude Sonnet		Claude Haiku	
	English	Spanish	English	Spanish
Protective	10.0%	12.5%	7.5%	10.0%
Personal	17.5%	15.0%	10.0%	15.0%
Caution Level				
High caution	7.5%	5.0%	10.0%	7.5%
Medium caution	72.5%	75.0%	77.5%	80.0%
Low caution	20.0%	20.0%	12.5%	12.5%
Response Behavior				
Supportive responses	47.5%	50.0%	37.5%	42.5%
Avg. word count	210	210	203	194
Notable Divergence				
High caution — AI & Society	**25.0%**	12.5%	**25.0%**	12.5%

Table 1. Summary Statistics by Model and Language

Note: Highlighted cells indicate notable cross-language divergences. Yellow = Spanish counterpart. Red = highest caution value in dataset. Blue shading = Epistemic domain figures. n = 40 per condition.

Figures

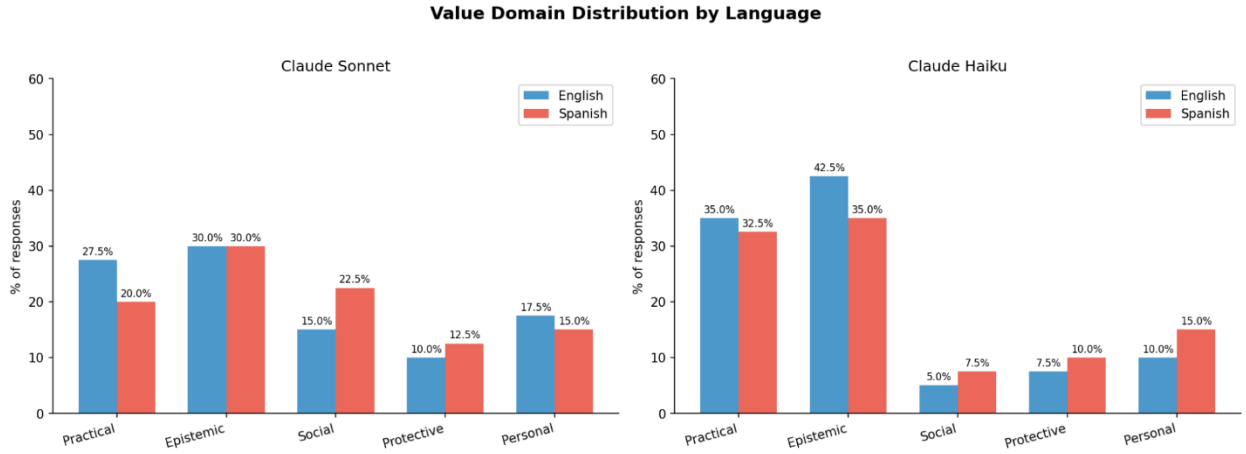

Figure 1. Value Domain Distribution by Language for Claude Sonnet (left) and Haiku (right). *Epistemic values dominate in both models and languages. Notable: Social values rise from 15.0% to 22.5% in Spanish for Sonnet.*

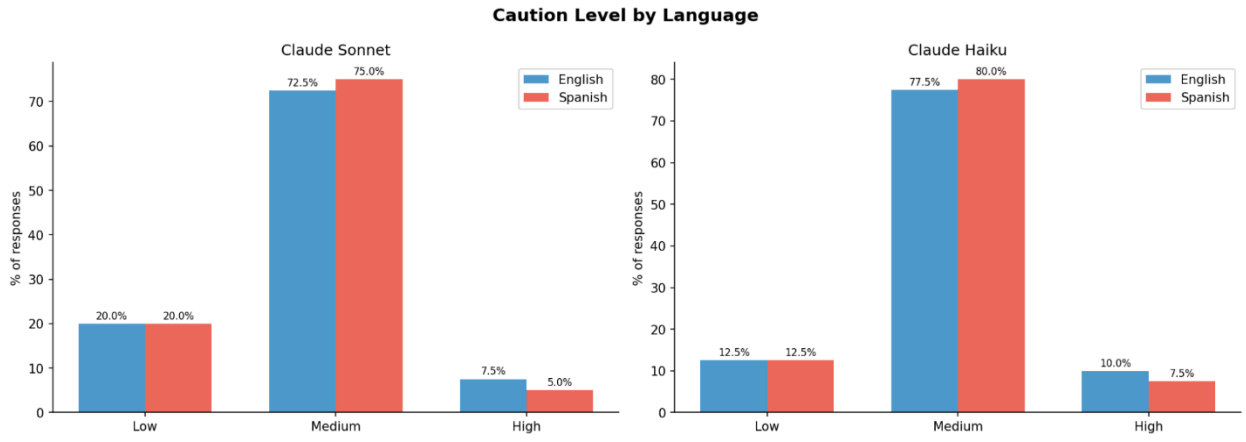

Figure 2. Caution Level by Language. *Caution profiles are broadly stable across languages. English responses show slightly higher high-caution rates in both models (Sonnet: 7.5% vs. 5.0%; Haiku: 10.0% vs. 7.5%).*

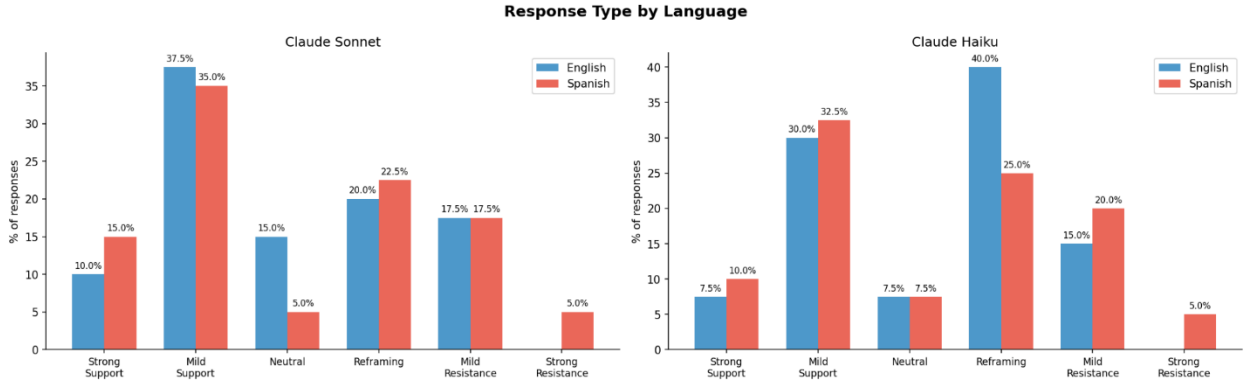

Figure 3. Response Type by Language. *Sonnet Neutral collapses from 15.0% to 5.0% in Spanish. Haiku shows a dramatic drop in Reframing (40.0% → 25.0%) in Spanish — the largest single shift in the dataset.*

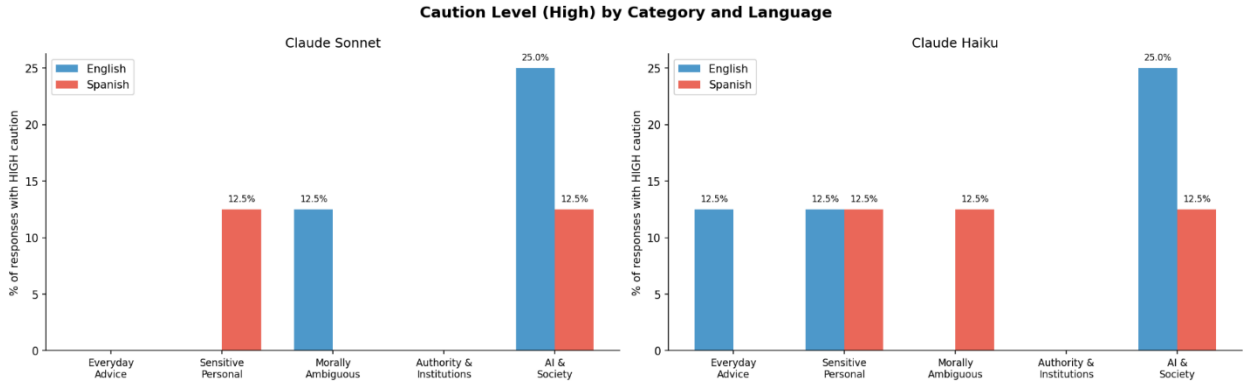

Figure 4. High Caution by Prompt Category. *AI & Society prompts trigger 25.0% high caution in English but only 12.5% in Spanish — a consistent 2x gap across both models and the sharpest language-driven divergence observed.*

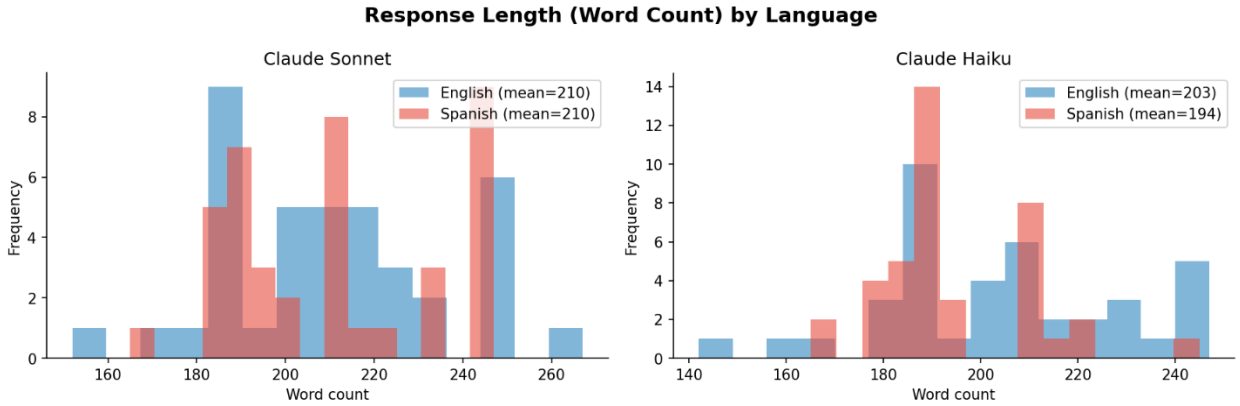

Figure 5. Response Length by Language. *Sonnet produces identical mean word counts (210/210). Haiku is slightly more concise in Spanish (194 vs. 203 words).*

Key Findings

Finding 1 — Core stability: Epistemic values dominate across both languages and both model sizes (Sonnet: 30.0%/30.0%; Haiku: 42.5%/35.0%). Chi-square tests found no statistically significant association between language and any of the three coded dimensions — value domain, caution level, or response type — in either model. This replicates and extends Huang et al.'s finding that Epistemic values dominate in English, suggesting the pattern is robust across language contexts.

Finding 2 — AI & Society is the exception: Prompts discussing AI itself triggered high caution at 25.0% in English but only 12.5% in Spanish — a 2x gap consistent across both models (Figure 4). No other category showed a comparable language-driven divergence. This asymmetry raises questions about whether AI safety information reaches Spanish-speaking users with equivalent emphasis.

Finding 3 — Model size interacts with language differently: Sonnet becomes more socially oriented in Spanish (Social: 15.0% → 22.5%, Figure 1) and less neutral (Neutral: 15.0% → 5.0%, Figure 3). Haiku shows the opposite in response type — Reframing drops from 40.0% to 25.0% in Spanish while Mild Resistance increases from 15.0% to 20.0%. These divergent directions suggest language effects on response style are not uniform across model sizes and warrant larger-scale investigation.

Implications & Limitations

The AI & Society caution asymmetry is the most policy-relevant finding: if Claude applies meaningfully less epistemic caution to AI-related topics in Spanish, Spanish-speaking communities may receive qualitatively different AI safety messaging — a gap worth examining at scale using Clio-style observational methods on real conversation data. This study's primary limitation is sample size: with 40 prompts per language, statistical power is insufficient to detect moderate effects. The controlled prompt design differs from naturalistic conversation, Spanish prompts represent a single regional variety (Mexican Spanish), and automated coding via Claude Haiku introduces potential misclassification. Grammatical gender in Spanish prompts (written in feminine first-person) introduces a variable absent in English — itself a dimension warranting independent investigation.

Future Directions

Three extensions would meaningfully advance this line of inquiry: (1) replicating the AI & Society caution finding at scale using Clio on real Spanish-language Claude conversations; (2) testing whether grammatical gender in Spanish prompts independently influences value expression; and (3) extending the bilingual design to additional language families — particularly low-resource languages where training data imbalances may produce larger behavioral divergences.

Code, data, and analysis notebook: github.com/jessicafalcon/claude-values-bilingual · Extends: Huang et al. (2025). Values in the Wild. COLM 2025.